

Jordan Lee

+1 (555) 010-2468 | jordan.lee@example.com | example.com/profile/jordan-lee | example.com/code/jordan-lee

EDUCATION

Western State University M.S. in Computer Engineering	Metro City, USA Sep 20XX - Expected Jun 20XX
Eastern Institute of Technology B.S. in Electrical Engineering	Metro City, Country Sep 20XX - Jun 20XX

EXPERIENCE

Mobility Systems Company Machine Learning Engineering Intern	Metro City, Country Oct 20XX - May 20XX
---	--

- Built an industrial-diagnostics branch of a larger inspection system, automating triage of 800+ sensor signals per case through ML-extracted feature schemas, evidence-grounded review, and explicit abstention on low-confidence findings; after launch, the overall system reduced its pending-case backlog by 68%.
- Designed a role-aware routing and assembly layer that restricted each of 3 specialist agents and 1 independent reviewer to its in-scope signals, keeping findings traceable to source data and preventing cross-agent hypothesis contamination.
- Improved claim correctness by 7%, evidence-citation quality by 9%, and tool-use accuracy by 8% over the base model on held-out diagnostic cases by fine-tuning a domain-specialist adapter on validated tool-use trajectories with assistant-only loss masking.
- Applied GRPO with grouped multi-step tool-use rollouts and a composite reward to raise mean diagnostic accuracy 8% and cut end-to-end latency 5% over the SFT baseline, with safeguards against fabricated citations, redundant tool calls, and claim under-generation.

Cloud Systems Capstone Machine Learning Engineer, Edge AI Program	Metro City, USA Dec 20XX - Jun 20XX
--	--

- Enabled quantization-aware quality recovery through LoRA distillation by profiling layer sensitivity with KL divergence, selecting low-sensitivity target layers, and inserting fake-quantization Q/DQ nodes before TensorRT conversion.
- Reduced mean generated-image MSE by 24% and cross-sample MSE variance by 27% versus the pre-distillation checkpoint through target-layer LoRA distillation from a multi-step teacher trajectory on a 90K-image public dataset.
- Reduced peak VRAM by 68% (8,400 to 2,700 MB) and inference latency by 66% (3,200 to 1,100 ms) versus FP16 PyTorch on an embedded NVIDIA platform, outperforming direct TensorRT INT8 by 38% and 18%, respectively, while limiting quantization MSE to 0.041 through an explicit Q/DQ engine.

PROJECTS

Agent Runtime Suite Owner TypeScript, Agent Runtime, Multi-Agent Systems example.com/code/agent-runtime	Aug 20XX - Present
--	--------------------

- Improved planted-defect localization from 82% to 94% by replacing fixed-perspective analysis with an agent-as-tool architecture that dynamically dispatched tasks to role-specialized sub-agents with isolated prompts, toolsets, and contexts.
- Hardened a multi-agent fan-out/fan-in runtime against nested-pool deadlocks, reuse of incomplete streamed responses, and tool-result loss during context compaction by separating concurrency pools, containing faults at the branch level, and gating cache writes on explicit stream completion.
- Kept model-bound working context below 10K tokens throughout a 100-turn stress test while cumulative raw conversation grew 100x, using independently budgeted context layers and staged compaction, with SQLite checkpoints and hybrid FTS5/vector retrieval for recoverable state and long-term memory.

Research-Agent Evaluation Framework Contributor LLM Evaluation, Agent Benchmarking example.com/code/research-evaluation	Feb 20XX - Jul 20XX
--	---------------------

- Built and upstreamed a unified evaluation infrastructure for an open-source research-agent framework, integrating 8 RAGAS/ALCE-inspired metrics spanning answer relevance, source quality, evidence faithfulness, and citation behavior.
- Validated evaluator fidelity with a Kendall rank correlation of 0.89 and an 8.5pp mean absolute deviation between injected and measured degradation across 400+ report-level trials by progressively removing citations, sources, sections, and factual claims in 10-percentage-point increments.
- Identified 3 structural defects in a multi-agent research pipeline across stability, sourcing, and parameter handling by combining layered instrumentation with systematic error attribution; mapped each failure to the responsible module to guide targeted fixes.

SKILLS

- Programming:** Python, TypeScript, SQL, Bash, Git
- LLM & Agents:** PyTorch, LangGraph, agent orchestration, Pydantic, RAG, MCP, FastAPI, OpenAI SDK, agent evaluation
- Post-Training & Optimization:** LoRA/QLoRA, SFT, GRPO, TensorRT Quantization, knowledge distillation, fine-tuning