

Jordan Lee

+1 (555) 010-2468 | jordan.lee@example.com | example.com/code/jordan-lee

Date of birth: 14 Mar 1999 | Nationality: Canadian

EDUCATION

Western State University | M.S. in Computer Engineering | Metro City, USA | Sep 2024 - Expected Jun 2026

Eastern Institute of Technology | B.S. in Electrical Engineering | Metro City, Country | Sep 2018 - Jun 2022

EXPERIENCE

Lakeside Auto Parts | Mechanical Design Engineer | Metro City, Country | Jun 2025 - Present

- Designed stamping dies for automotive brackets in SolidWorks, cutting the scrap rate from 6% to 4% on two press lines.
- Ran tolerance stack-up analyses for 12 production parts and signed off first-article inspections with the supplier.

Mobility Systems Company | Machine Learning Engineering Intern | Metro City, USA | Oct 2024 - May 2025

- Improved diagnostic accuracy by 35% after fine-tuning a domain adapter on validated tool-use trajectories with assistant-only loss masking.
- Designed a routing layer that limits each of 3 specialist agents and an independent reviewer to their in-scope signals, cutting reviewer disagreement with specialist findings from 14% to 6%.
- Trained the triage agent with GRPO on grouped tool-use rollouts and a reward that penalises redundant calls, cutting tool calls per case 18% and end-to-end latency 5% versus the SFT baseline at equal accuracy.
- Wrote the evaluation harness the team used to compare 14 adapter checkpoints on accuracy, citation quality and latency, catching 2 accuracy regressions before release.
- Fine-tuned the adapter with assistant-only loss masking so the model would learn to reproduce the tool outputs more faithfully.
- Built a diagnostics triage branch for an industrial inspection system that screens 800+ sensor signals per case with ML-extracted features, cutting the pending-case backlog 68% in the eight weeks after launch.

Eastern Robotics Co. | Junior Software Engineer | Metro City, Country | Aug 2022 - Jul 2024

- Cut GPU memory for fine-tuning the perception models by 4x by switching from FP32 to BF16 mixed precision.
- Reduced p95 API latency from 420 ms to 180 ms by adding a request cache and batching sensor reads, with load tests that fail the build if p95 exceeds 200 ms.
- Maintained the CI pipeline for the perception team's model releases and adds automated regression checks that shortened release cycles to 3 days.
- Migrated 30 robot-fleet services from cron jobs to an event queue with retries and dead-letter handling, removing the nightly backlogs that delayed morning dispatch.

PROJECTS

Agent Runtime Suite | Owner | TypeScript, Multi-Agent Systems | Aug 2025 - Present

- Drove adoption of AI-first engineering practices across the platform, accelerating delivery and improving outcomes for downstream teams.
- Kept working context under 10K tokens across a 100-turn stress test while the raw conversation grew 100x, using budgeted context layers and staged compaction.
- Separated concurrency pools and gated cache writes on stream completion, removing nested-pool deadlocks and lost tool results under 50-way fan-out.

Research-Agent Evaluation Framework | Contributor | LLM Evaluation | Feb 2025 - Jul 2025

- Integrated 8 citation and faithfulness metrics into an open-source research-agent framework's evaluation module.
- Showed the evaluator tracks injected degradation with a Kendall correlation of 0.89 across 400+ report-level trials that removed citations, sources and claims.
- Cut the pending-case backlog by two-thirds with a triage branch that screens 800+ sensor signals per case using ML-extracted features.

SKILLS

Programming: Python, TypeScript, Git

ML & Agents: PyTorch, LoRA, GRPO, agent evaluation