

Morgan Patel

+1 (555) 010-4410 | morgan.patel@example.com | example.com/code/mpatel

Date of birth: 21 Jul 1998 | Nationality: Indian

EDUCATION

Ridgeway University | Ph.D. candidate in Statistics | Metro City, USA | Sep 2021 - Expected May 2026

Ridgeway University | B.S. in Mathematics | Metro City, USA | Sep 2016 - May 2020

EXPERIENCE

Ridgeway University | Research Assistant, Statistical Learning Lab | Metro City, USA | Jun 2020 - Aug 2021

- Owned the lab's simulation pipeline on the shared cluster and the reproducibility checks other students relied on.
- Derived a variance bound for a sparse regression estimator that tightens the previous bound by a log factor; it is now Section 3 of a paper under review at JASA.
- Taught weekly recitations for 60 students in graduate probability, writing 12 problem sets and earning a 4.8/5 teaching rating.
- Released an open-source R package for high-dimensional covariance estimation while maintaining the lab's shared cluster, organizing the weekly reading group and grading for two courses, which was downloaded 3,000 times in its first year.

Northpeak Capital | Quantitative Research Intern | Metro City, USA | Jun 2025 - Aug 2025

- Built a short-horizon order-book imbalance signal for liquid index futures that raised the desk book's Sharpe ratio from 1.1 to 1.5 over 18 months of out-of-sample backtest after costs.
- Cut the signal's daily turnover from 34% to 21% with a cost-aware position smoother, keeping 90% of gross returns and lowering estimated slippage by a third.
- Confirmed the forecast gain over the nested HAR-RV baseline with standard Diebold-Mariano tests across the 30 indices.
- Joining 120 microstructure features point-in-time across six venues, deduplicating late prints and versioning each schema, built a feature store the team reused in two later projects.
- Annualized the signal's daily Sharpe ratio by multiplying it by 252 before reporting it to the desk.
- Documented the backtest assumptions, transaction-cost model and known failure regimes in the desk's research wiki, which the next intern cohort used to onboard in their first week.

PROJECTS

Volatility Forecasting Study | Independent Research | Python, PyTorch | Jan 2024 - Present

- Beat a HAR-RV baseline's out-of-sample QLIKE loss by 7% on 30 equity indices with a temporal convolutional model trained on realized-volatility features.
- Cut forecast error from 0.20 to 0.15, a 33% improvement, by adding realized-volatility features and an asymmetric loss.
- Improved the model's directional hit rate by 6% on 30 equity indices, from 52% to 58%, with a temporal convolutional model.

Kaggle Market Prediction Competition | Team of 3 | Python | Mar 2023 - Jun 2023

- Placed 41st of 2,900 teams with a gradient-boosting ensemble on 300 engineered features, finishing in the top 2% of the private leaderboard.
- Cut validation leakage by switching to time-grouped folds, which closed a 0.02 gap between local validation and leaderboard scores.
- Wrote the team's feature-selection script using permutation importance, trimming 900 candidate features to 300 without losing validation score.

SKILLS

Programming: Python, R, PyTorch

Methods: time-series econometrics, high-dimensional statistics, gradient boosting