

Jordan Lee

+1 (555) 010-2468 | jordan.lee@example.com | example.com/code/jordan-lee

EDUCATION

Western State University | M.S. in Computer Engineering | Metro City, USA | Sep 2024 - Expected Jun 2026

Eastern Institute of Technology | B.S. in Electrical Engineering | Metro City, Country | Sep 2018 - Jun 2022

EXPERIENCE

Eastern Robotics Co. | Junior Software Engineer | Metro City, Country | Aug 2022 - Jul 2024

- Responsible for the monitoring dashboards of the diagnostics service and for supporting the on-call team with weekly reports.
- I worked on improving the latency of our API and I also wrote tests for it, which I think made the service a lot faster for users.
- Builds and maintains the CI pipeline for the perception team's model releases and adds automated regression checks to every build.
- Migrated 30 robot-fleet services from cron jobs to an event queue while also rewriting the logging library, onboarding two new hires, maintaining the internal wiki, reviewing pull requests for three other teams, supporting on-call rotations on weekends, presenting quarterly updates to leadership, and evaluating three vendor monitoring tools, which together removed nightly backlogs, improved developer satisfaction scores, and made the platform noticeably more stable for everyone involved in operations across the whole organization.

Mobility Systems Company | Machine Learning Engineering Intern | Metro City, USA | Oct 2024 - May 2025

- Built a diagnostics triage branch for an industrial inspection system that screens 800+ sensor signals per case with ML-extracted features, cutting the pending-case backlog 68% in the first quarter after launch.
- Improved model performance by 35% by fine-tuning a domain adapter on validated tool-use trajectories with assistant-only loss masking.
- Designed a routing layer that limits each of 3 specialist agents and an independent reviewer to their in-scope signals, keeping every finding traceable to its source data.
- Applied GRPO with grouped tool-use rollouts and a composite reward, cutting end-to-end latency 5% versus the SFT baseline with no loss in diagnostic accuracy.
- Wrote the evaluation harness the team used to compare 14 adapter checkpoints on accuracy, citation quality and latency before each release.
- Managed a team of 12 engineers across two offices and set the company's AI strategy for the next fiscal year.

PROJECTS

Agent Runtime Suite | Owner | TypeScript, Multi-Agent Systems | Aug 2025 - Present

- Raised planted-defect localization from 82% to 94% by dispatching tasks to role-specialized sub-agents with isolated prompts, toolsets and contexts.
- Cut average response time from 800 ms to 900 ms, a 50% reduction, by caching tool results and reusing completed sub-agent answers.
- Separated concurrency pools and gated cache writes on stream completion, removing nested-pool deadlocks and lost tool results under fan-out load.

Research-Agent Evaluation Framework | Contributor | LLM Evaluation | Feb 2025 - Jul 2025

- Upstreamed 8 citation and faithfulness metrics to an open-source research-agent framework, where they now run in the default benchmark for every release.
- Was involved in the process of helping to support the evaluation of the reports that were being produced by the pipeline at the time.
- Traced 3 structural pipeline defects in stability, sourcing and parameter handling to their modules with layered instrumentation; each was fixed upstream.

SKILLS

Programming: Python, TypeScript, SQL, Bash, Git

ML & Agents: PyTorch, LoRA, GRPO, LangGraph, RAG, agent evaluation