

Jordan Lee

+1 (555) 010-2468 | jordan.lee@example.com | example.com/code/jordan-lee

Date of birth: 14 Mar 1999 | Nationality: Canadian

EDUCATION

Western State University | M.S. in Computer Engineering | Metro City, USA | Sep 2024 - Expected Jun 2026

Eastern Institute of Technology | B.S. in Electrical Engineering | Metro City, Country | Sep 2018 - Jun 2022

EXPERIENCE

Eastern Robotics Co. | Junior Software Engineer | Metro City, Country | Aug 2022 - Jul 2024

- Owned the diagnostics service's monitoring dashboards across two major releases and the on-call rotation that used them.
- Reduced p95 API latency from 420 ms to 180 ms; my changes added a request cache and batched sensor reads, with load tests to keep it there.
- Maintained the CI pipeline for the perception team's model releases and adds automated regression checks that shortened release cycles to 3 days.
- Migrated 30 robot-fleet services from cron jobs to an event queue with retries and dead-letter handling, removing the nightly backlogs that delayed morning dispatch.

Mobility Systems Company | Machine Learning Engineering Intern | Metro City, USA | Oct 2024 - May 2025

- Built a diagnostics triage branch for an industrial inspection system that screens 800+ sensor signals per case with ML-extracted features, cutting the pending-case backlog 68% in the first quarter after launch.
- Improved diagnostic accuracy by 35% after fine-tuning a domain adapter on validated tool-use trajectories with assistant-only loss masking.
- Designed a routing layer that limits each of 3 specialist agents and an independent reviewer to their in-scope signals, keeping every finding traceable to its source data.
- Applied GRPO with grouped tool-use rollouts and a composite reward, cutting end-to-end latency 5% versus the SFT baseline with no loss in diagnostic accuracy.
- Wrote the evaluation harness the team used to compare 14 adapter checkpoints on accuracy, citation quality and latency before each release.
- Led a team of 4 engineers to design and ship the triage branch's escalation workflow for the on-call reviewers.

PROJECTS

Agent Runtime Suite | Owner | TypeScript, Multi-Agent Systems | Aug 2025 - Present

- Raised planted-defect localization from 82% to 94% by dispatching tasks to role-specialized sub-agents with isolated prompts, toolsets and contexts.
- Kept working context under 10K tokens across a 100-turn stress test while the raw conversation grew 100x, using budgeted context layers and staged compaction.
- Moved the runtime's evaluation jobs onto the KRT deployment stack, cutting the time to rerun a full benchmark from 40 to 12 minutes.

Research-Agent Evaluation Framework | Contributor | LLM Evaluation | Feb 2025 - Jul 2025

- Supported cross-functional teams on evaluation initiatives for the open-source framework, improving overall quality and consistency.
- Showed the evaluator tracks injected degradation with a Kendall correlation of 0.89 across 400+ report-level trials that removed citations, sources and claims.
- Traced 3 structural pipeline defects in stability, sourcing and parameter handling to their modules with layered instrumentation; each was fixed upstream.

SKILLS

Programming: Python, TypeScript, SQL, Bash, Git

ML & Agents: PyTorch, LoRA, GRPO, LangGraph, RAG, agent evaluation